

AI-AUTHORITY-001 — AI Authority Boundary

Status: FROZEN — MIP-001 W0 / W9 **Constitutional basis:** Articles 8, 20, 21, 27 **Executable form:** `pkg/ai/gateway`

1. The rule

AI has no evidence authority.

AI may propose, summarize, hypothesize, correlate, and recommend. It may not create, alter, qualify, or sign authoritative evidence. This holds for Aureum, God of EYS, Claude, GPT, in-house ML, and any future model — the boundary is a property of the architecture, not a policy about a particular vendor.

2. What AI may and may not do

May:

Reasoning · Discovery · Hypothesis generation · Investigation planning
Evidence summarization · Pattern analysis · Next-evidence recommendation
Natural-language investigation

May not:

alter evidence · delete evidence · change trust · change policy
approve disclosure · confirm privilege · suppress contradiction
qualify authoritative evidence · sign final qualification
determine legal liability · instruct a connector

The last is easy to overlook and the most dangerous: an AI that can direct an acquisition connector has acquired evidence authority indirectly, by choosing what enters the fabric. IEAP-001 §3 denies it explicitly.

3. The gateway is mandatory

No AI touches raw evidence. Every access passes `pkg/ai/gateway`:

AI Request → Purpose Check → Evidence Classification → Privilege Check
→ Protective Order → License / Rights → AI Processing Policy
→ Minimum Necessary Disclosure → Approved Evidence Projection → AI

RAW EVIDENCE —X— AI cannot access directly

APPROVED EVIDENCE PROJECTION —→ AI

The projection is the only surface AI sees. This matters commercially as much as legally: it is what prevents a customer's licensed data or another tenant's evidence from entering a model context.

4. Ten evaluation dimensions

Every access evaluates: purpose, classification, privilege, protective order, rights, license, jurisdiction, AI processing policy, recipient, minimum necessary disclosure.

Redacted evidence does not automatically become AI-safe (Article 21). Redaction serves a disclosure purpose; AI processing is a *different* purpose with different rights. A derivative cleared for opposing counsel's eyes is not thereby cleared for a training corpus.

5. AI contribution record (W9.4)

Every material AI contribution records:

Model · Model Version · Prompt Hash · Input Evidence IDs
Input Evidence Version IDs · Policy Version · Tool Calls
Output Hash · Human Reviewer · Purpose · Timestamp

This delivers Article 27 — **No Silent AI Influence**. If a model shaped a finding, the fact pack shows which model, on which evidence versions, under which policy, reviewed by whom.

Note **Input Evidence Version IDs**, not just IDs: a contribution made against version 2 of an artifact is not the same contribution if version 3 later exists.

6. Purpose separation (Article 20)

VIEW ≠ SEARCH ≠ COPY ≠ PRINT ≠ DOWNLOAD ≠ EXPORT
≠ REDISTRIBUTE ≠ AI_PROCESS ≠ RAG ≠ TRAIN

Ten separately-granted rights. **AI_PROCESS**, **RAG**, and **TRAIN** are three distinct grants: permission to run a model over evidence once is not permission to index it for retrieval, and neither is permission to train on it. Collapsing them is the single most common way evidence rights are silently exceeded.

7. Privileged material defaults

Unless explicitly authorized, privileged material gets:

No general search · No general RAG · No cross-case learning
No Common Fact Pack · No export

Default-deny, not default-allow-with-audit.

8. Where Aureum and God of EYS sit

Both are **Intelligence Plane** capabilities. Their outputs enter EQF as hypotheses or observations *requiring qualification* — never as findings.

Aureum / God of EYS → Hypothesis → EQF → Qualification

God of EYS may report *"these entities may be connected."* VERIQO then asks: what evidence supports this? what contradicts it? what alternative exists? what source is independent? what is still missing? That interrogation is what keeps analytics epistemically disciplined rather than merely fluent.

9. Implementation status

Component	Status
AI evidence gateway (10 checks, fail-closed)	IMPLEMENTED, UNIT_TESTED, ADVERSARIAL_TESTED
AI contribution record	IMPLEMENTED, UNIT_TESTED
Purpose-separated rights (10 rights)	IMPLEMENTED, UNIT_TESTED
Privileged default-deny	IMPLEMENTED, ADVERSARIAL_TESTED
Aureum AI	NOT BUILT — requires a model deployment
God of EYS	NOT BUILT — requires an investigation corpus
Connector-instruction denial	IMPLEMENTED as a constitutional check; no live connector exists to attack

The gateway is built and enforced **before** the intelligences that must pass through it exist. That ordering is deliberate: a boundary retrofitted after the systems it constrains have shipped is a boundary that will have been bypassed.